

Course Information



- **Instructor:** Ekaterina Vladimirovna Bolgova
 - **E-mail:** *ekaterina_bolgova@itmo.ru*
- **Course language:** English.
- **Classes starts at** September 14, 2019.
- **Duration:** 1 semester.
- **Classes' types:**
 - Lectures (4): 14/09, 21/09, 28/09, 05/10
 - Labs (programming tasks): starts at October 19, 2019
- **Exam (2/1/0 question(s)):**
 - Submit all labs (code + report)
 - Make presentation (at least 1)
 - Attendance & activity
 - Bonus: if make workshop





Basics of Parallel Programming with OpenMP

PhD Katerina Bolgova
eScience Research Institute & HPC Department

Background

- What are Parallel Computer Memory Architectures?


Distributed Memory

Shared Memory


Hybrid Shared-Distributed Memory

UMA

NUMA

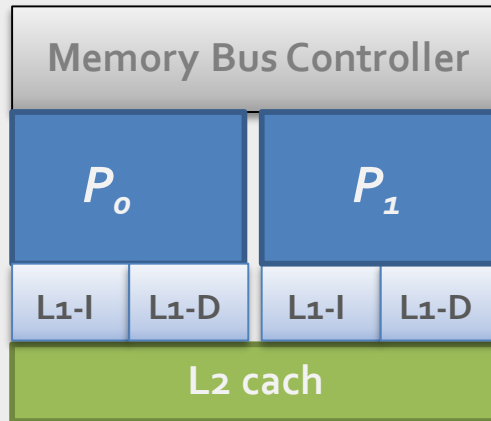


Symmetric Multiprocessor

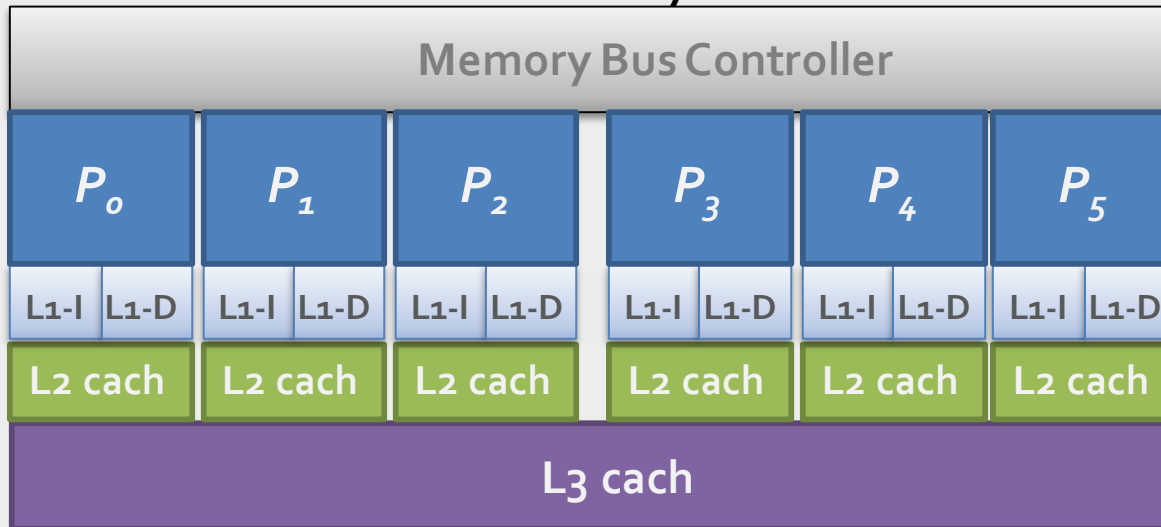


Background

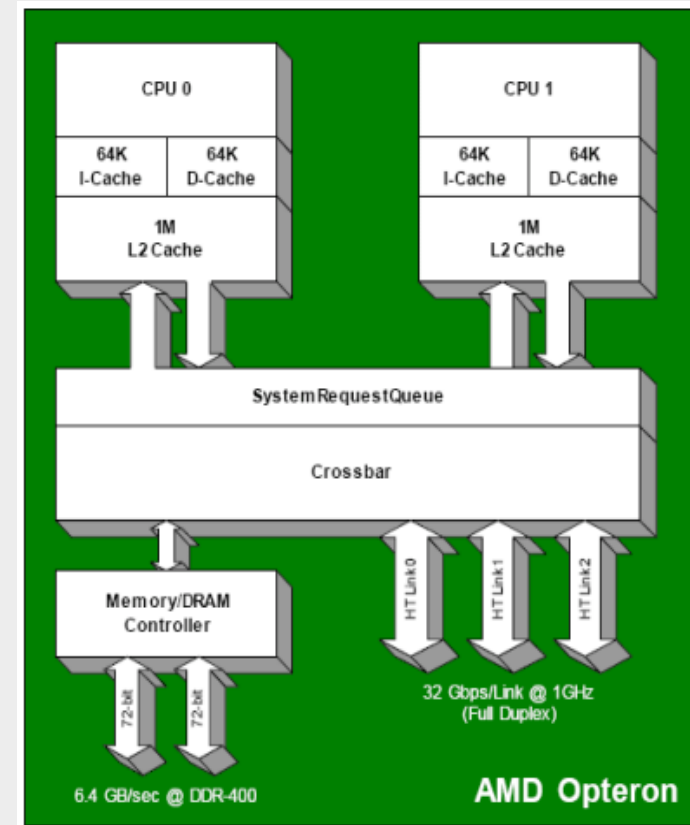
Intel Core 2 Duo



Intel Core i7



AMD Opteron



OpenMP Overview



OpenMP :

- is an Application Program Interface (**API**) that may be used to explicitly direct *multi-threaded, shared memory parallelism*
- is managed by the OpenMP Architecture Review Board (or *OpenMP ARB*)
- provides a portable, scalable model for developers of shared memory parallel applications.
- supports C/C++ and Fortran
- *is a set of compiler directives and library routines for parallel application programmers*

Brief history

- The OpenMP standard specification started in 1997. (The version for Fortran appeared).
- In 1998 the C/C++ standard was released
- Since 2005, C and Fortran specifications have been released together
- The latest version (4.0) of the specification was released in July 2013




Advantages of OpenMP

- Provide a *standard* among a variety of shared memory architectures
- Establish a simple and limited set of directives for programming shared memory machines
- Significant parallelism can be implemented by using just 3 or 4 directives
- Provide capability to *incrementally* parallelize a serial program
- The API is specified for C/C++ and Fortran
- Public forum for API and membership
- Most major platforms have been implemented including Unix/Linux platforms and Windows



OpenMP Programming Model

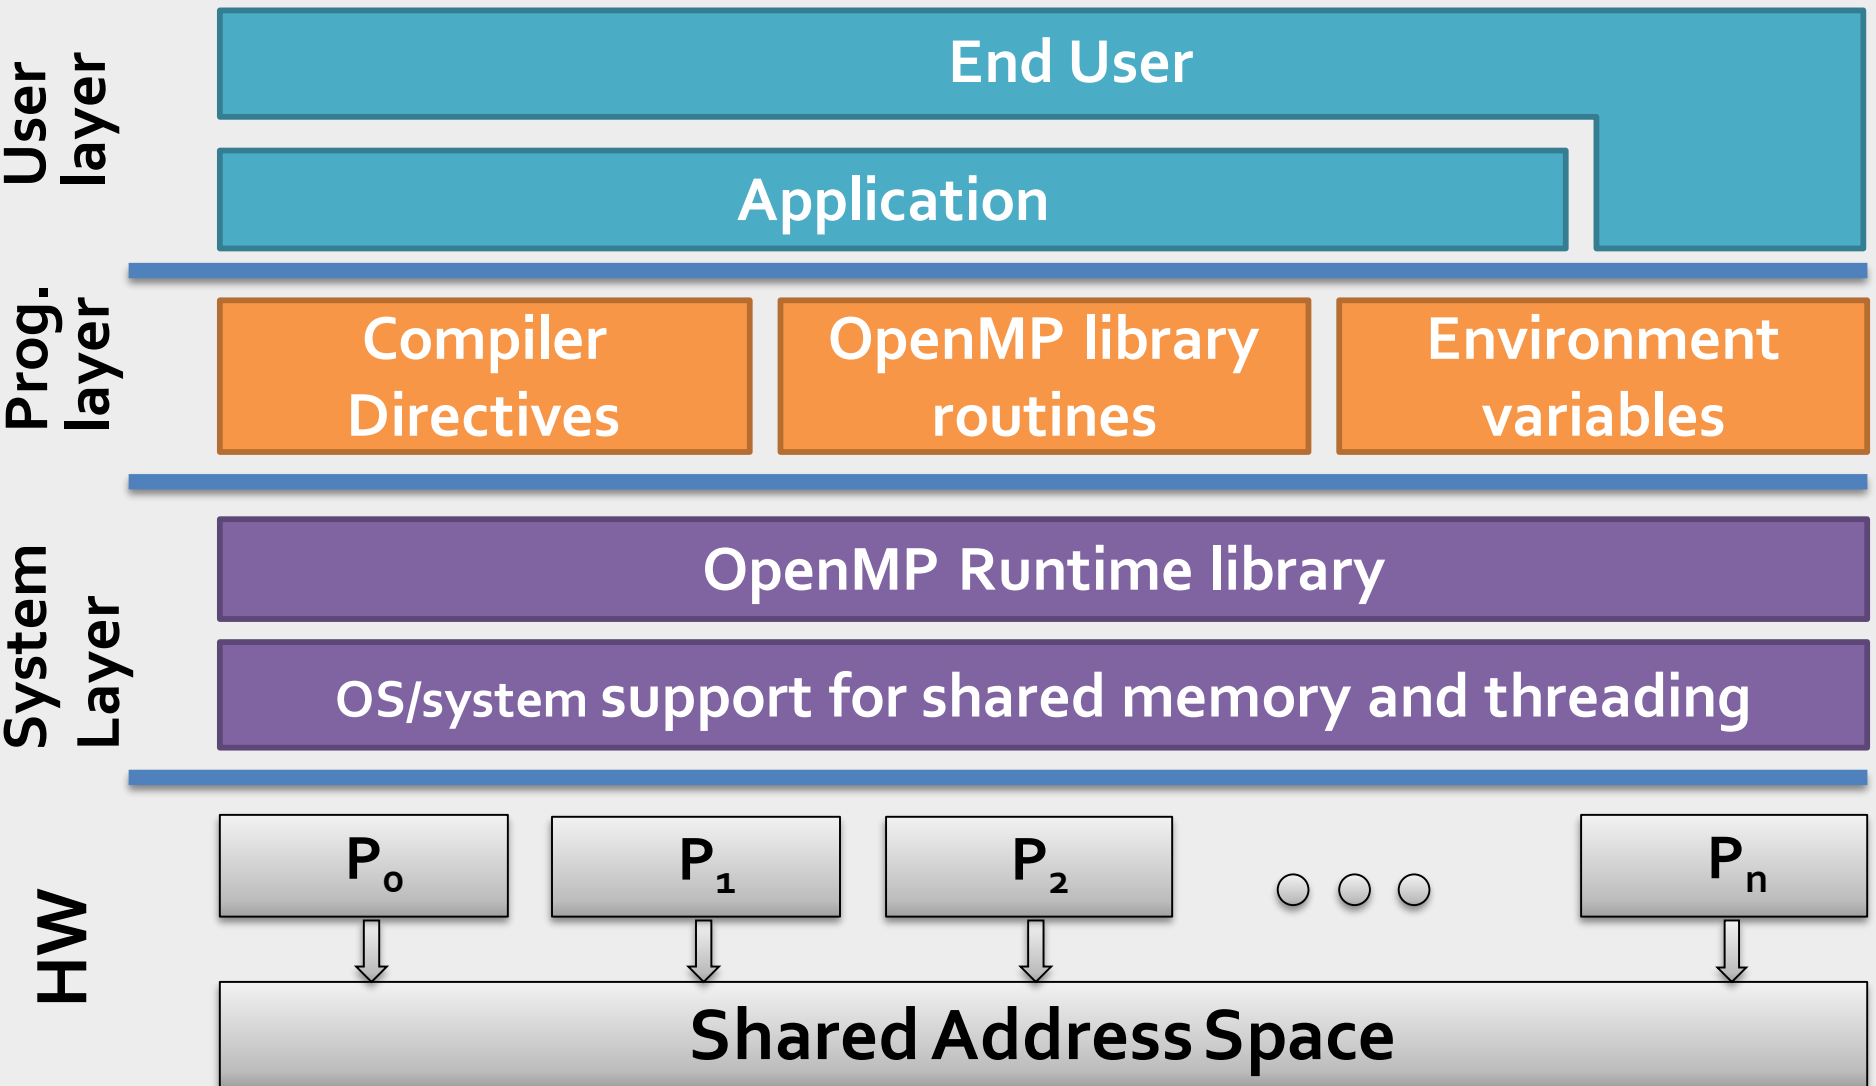


- OpenMP program begins as a single process: the *master thread*
- **FORK**: the master thread then creates a *team* of parallel *threads*
- Parallel execution
- **JOIN**: the team threads terminate, leaving only the master thread
- Parallelism added *incrementally*
- The number of parallel regions and the threads that comprise them are arbitrary
- Each thread can spawn another team of threads (nested parallelism)

FORK-JOIN Model



OpenMP Solution Stack



Source: Tim Mattson, Introduction to OpenMP



OpenMP Directives Format

```
#pragma omp directive-name [clause[ [,] clause]...] new-line
```

- Each directive starts with **#pragma omp**
- Compilers can ignore OpenMP directives and conditionally compiled code if support of the OpenMP API is not provided
- Directives are case-sensitive
- An OpenMP executable directive applies to at most one succeeding statement, which must be a structured block
- Long directive lines can be "continued" on succeeding lines with a backslash ("\") at the end of a directive line
- *for C/C++ it needs to include the `<omp.h>` header file*

Example:

```
#pragma omp parallel default(shared) \  
                    private(beta,pi)
```



Thread Creation

```
#pragma omp parallel [clause ...] newline  
    structured_block
```

- When thread reaches this directive, it creates a *team* of threads
- All threads will execute code of this parallel region
- There is an implied barrier at the end of a parallel section
- The number of threads in the team remains constant for the duration of that parallel region

Example:

```
#include <omp.h>
int main(){
    #pragma omp parallel
    {
        int ID = omp_get_thread_num();
        cout<<" hello ("<< ID <<") \n";
        cout<<" world ("<< ID <<") \n";
    }
}
```

Annotations:

- OpenMP include file (points to `<omp.h>`)
- Parallel region (points to `#pragma omp parallel`)
- End of the Parallel region (points to the closing brace `}`)



Thread Creation

Sample Output:



```
hello <1>
hello <2>
hello <5> hello <7> hello <10>
world <10>
hello <6> hello <9> hello <11>
world <11>

world <5>
world <9> world <1> hello <4>

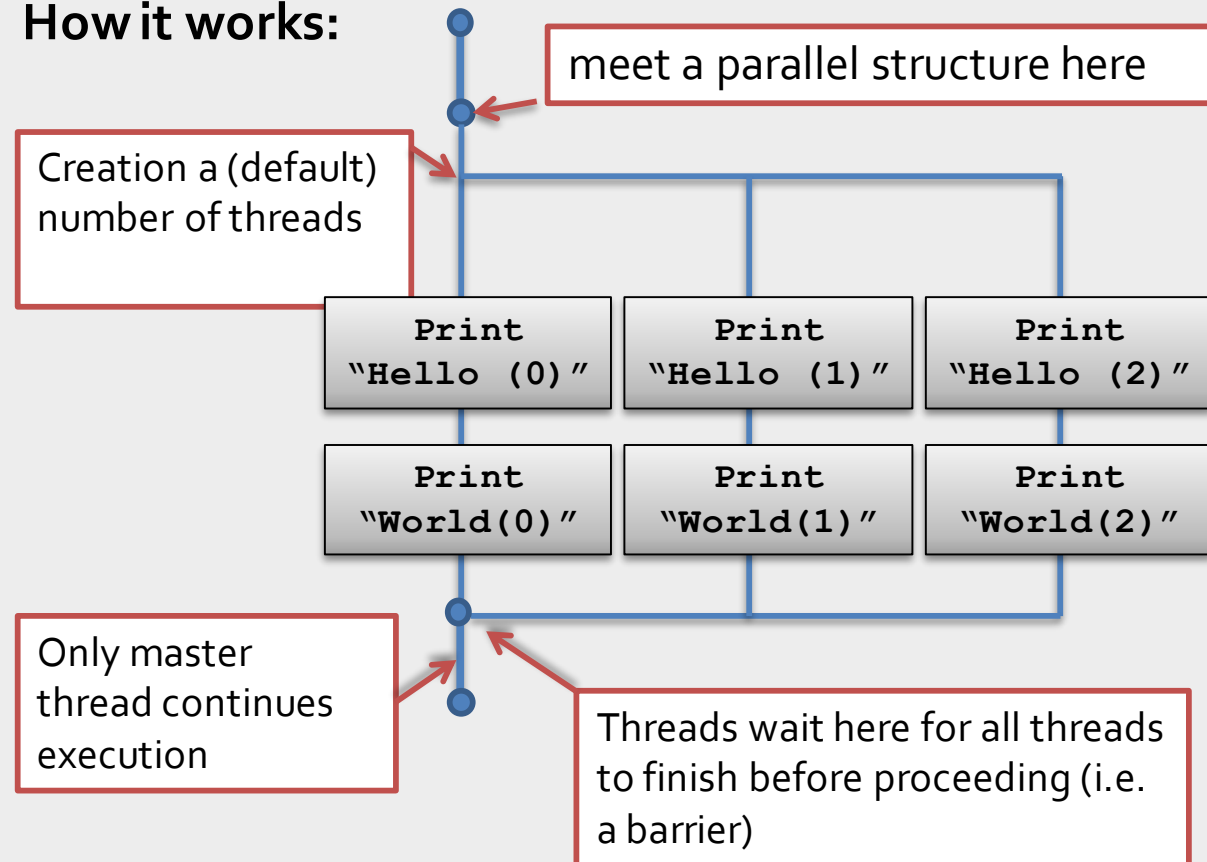
world <7>

hello <3>

world <6>
hello <8>
world <8>
hello <0>
world <2>

world <0>
world <3>
world <4>
```

How it works:



Clauses of Parallel Construct Overview

- **num_threads** (integer-expression) – explicitly determine the number of threads
- **private** (list) – each thread will have a local copy and use it as a temporary variable
- **firstprivate** (list) – like *private* except initialized to original value
- **shared** (list) – the data within a parallel region is shared
- **default (shared | none)** – allows the programmer to state that the default data scoping within a parallel region will be either *shared*, or *none*
- **reduction** (operator :list) – a safe way of joining work from all threads after construct
- **if** (scalar-expression) – This will cause the threads to parallelize the task only if a condition is met. Otherwise the block executes serially



Defining the Time of Program Execution

- Time in Seconds since a fixed point in the past:

```
double omp_get_wtime(void)
```

- Scheme of using the **omp_get_wtime** function:

```
double t_1, t_2, dt;  
t_1 = omp_get_wtime();  
...  
t_2 = omp_get_wtime();  
dt = t_2 - t_1;
```



Determining Number of Threads

- The number of threads in a parallel region is determined by the following factors, in order of precedence:

- Evaluation of the *if* clause

```
#pragma omp parallel if (NMAX>LIMIT)
```

- Setting of the *num_threads* clause

```
#pragma omp parallel num_threads(4)
```


- Use of the *omp_set_num_threads()* library function

```
void omp_set_num_threads (int num_threads);
```

- Setting of the *OMP_NUM_THREADS* environment variable



Sample: PI evaluation



We know that:

$$\int_0^1 \frac{4}{1+x^2} dx = \pi$$




We can approximate the integral as a sum of rectangles:

$$\sum_{i=0}^N F(x_i) \Delta x \approx \pi, \quad \text{where} \quad \begin{array}{l} \Delta x \text{ -- rectangle width} \\ F(x_i) \text{ -- rectangle height} \end{array}$$



PI: Serial Program

```
static long cntSteps=10000000;  
double step;  
  
void main() {  
    double x, pi, sum = 0.0;  
    int i;  
  
    step = 1.0/(double)cntSteps;  
  
    for (i=0; i<cntSteps; i++)    {  
        x = (i + 0.5)*step;  
        sum = sum + 4.0/(1.0+ x*x) ;  
    }  
    pi = sum*step;  
  
    cout << "PI = "<<setw(15)<<setprecision(12)<<pi;  
}
```



PI: Simple Parallel Program

```
#include <omp.h>
static long cntSteps=10000000;      double step;
#define thrdsCount = 2              The SPMD design pattern
                                     (Single Program Multiple Data)
void main() {
    double pi, sum[thrdsCount];
    int i;
    step = 1.0/(double)cntSteps;
    #pragma omp parallel num_threads(thrdsCount)
    {
        int i,id; double x;
        id = omp_get_thread_num();
        for (i=0, sum[id]=0; i<cntSteps; i+=thrdsCount) {
            x = (i + 0.5)*step;
            sum[id] += 4/4100 (1.0/x)*x*x;
        }
    }
    for(i=0; i<thrdsCount; i++)
        pi += sum[i] * step;
    cout << "PI = " << setw(15) << setprecision(12) << pi;
}
```

PI: Simple Parallel Program

```
#include <omp.h>
static long cntSteps=10000000; double step;
#define thrdsCount = 2
void main() {
    double pi, sum[thrdsCount];
    int i;
    step = 1.0/(double)cntSteps;
    #pragma omp parallel num_threads(thrdsCount)
    {
        int i,id; double x;
        id = omp_get_thread_num();
        for (i=id, sum[id]=0.0;i<cntSteps; i+=thrdsCount) {
            x = (i + 0.5)*step;
            sum[id] += 4.0/(1.0 + x*x);
        }
    }


    for(i=0, pi=0.0; i<thrdsCount;i++)
        pi += sum[i] * step;
    cout << "PI = " << setw(15) << setprecision(12) << pi;
}
```

#Th	T_p	S_p
1	0,79	1
2	0,51	1,55
4	0,46	1,72
8	0,43	1,84
16	0,36	2,19



False sharing

- Modifying variables that reside on the same cache line by threads on different processors will cause the cache lines to “slosh back and forth” between threads.
- It’s the reason of “**false sharing**”
- Combining scalars to an array program is a very common way to support creation of an SPMD
- But most likely the array elements are contiguous in memory and hence share cache lines
- Simple solution: Pad arrays



PI: Simple Parallel Program

```
#include <omp.h>
static long cntSteps=10000000; double step;
#define thrdsCount = 2
#define PAD 8 // assume 64 bytes L1 cache line size
void main() {
    double pi, sum[thrdsCount][PAD];
    int i;
    step = 1.0/(double)cntSteps;
    #pragma omp parallel num_threads(thrdsCount)
    {
        int i, id, double x;
        id = omp_get_thread_num();
        for (i=id, sum[id][0]=0.0; i<cntSteps; i+=thrdsCount) {
            x = (1 + 0.5)*step;
            sum[id][0] += 4.0/(1.0 + x*x);
        }
    }

    for(i=0, pi=0.0; i<thrdsCount;i++)
        pi += sum[i][0] * step;
    cout << "PI = " << setw(15) << setprecision(12) << pi;
}
```

#Th	1 st T_p	2 nd T_p	S_p
1	0,79	0,78	1
2	0,51	0,41	1,9
4	0,46	0,23	3,39
8	0,43	0,14	5,57
16	0,36	0,13	6,0

Padding arrays doesn't resolve the "false sharing" problem

Work-Sharing in OpenMP

Types of Work-Sharing Constructs (WShCs):

- **LOOP CONSTRUCT**
- **SECTIONS/SECTION CONSTRUCTS**
- **SINGLE CONSTRUCT**
- **WORKSHARE CONSTRUCT**
- **TASK CONSTRUCT**


- A **WShC** divides the execution of code region among the members of the team
- **WShCs** do not launch new threads
- There is no implied barrier upon entry to a **WShC**, however there is an implied barrier at the end of a **WShC**.
- A **WShC** must be enclosed within a parallel region
- **WShCs** must be encountered by all members of a team or none at all



Loop Construct

Directive format: `#pragma omp for [clause,...] new-line
for-loop`

```
#pragma omp parallel
{
    #pragma omp for
    for (i=0;i<N;i++) {
        work_for_loop(i);
    }
}
```



- **for** directive specifies that the iterations of the loop (following it) must be executed in parallel by the team
- The loop variable (*i* in this case) is made “private” to each thread by default



Clauses of For Construct Overview

- **private** (list) – each thread will have a local copy and use it as a temporary variable
- **firstprivate** (list) – like *private* except initialized to original value
- **lastprivate** (list) – each variable from the sequentially last iteration of the associated loop is assigned to the variable's original object
- **reduction** (operator: list) – a safe way of joining work from all threads after construct
- **schedule** (type[, chunk]) – describes how iterations of the loop are divided among the threads in the team
- **collapse** (n) — Specifies how many loops in a nested loop should be collapsed into one large iteration space and divided according to the schedule clause
- **ordered** – the iterations of the loop must be executed as they would be in a serial program
- **nowait** – threads do not synchronize at the end of the parallel loop

1. Task for report



Construct Comparison

Serial code:

```
for(i=0;i<n;i++) {  
    a[i] = a[i-1]+c[i];  
}
```

Parallel region: SPMD pattern

```
#pragma omp parallel  
{  
    int id, i, Nths, istart, iend;  
    id = omp_get_thread_num();  
    Nths = omp_get_num_threads();  
    istart = id * N / Nthrds;  
    iend = (id+1) * N / Nthrds;  
    if (id == Nthrds-1)  
        iend = N;  
    for(i=istart;i<iend;i++) {  
        a[i] = a[i-1] + c[i];  
    }  
}
```

Parallel region: worksharing

```
#pragma omp parallel  
{  
    #pragma omp for  
    for(i=0;i<n;i++) {  
        a[i] = a[i-1] + c[i];  
    }  
}
```

or

```
#pragma omp parallel for  
for(i=0;i<n;i++) {  
    a[i] = a[i-1] + c[i];  
}
```



Sample of For construct

```
#include <omp.h>
#define NMAX 1000
main () {
    int i, j, sum;
    float a[NMAX][NMAX];
    <data initialization>

    #pragma omp parallel for shared(a) private(i,j,sum)
    {
        for (i=0; i < NMAX; i++) {
            sum = 0;
            for (j=0; j < NMAX; j++)
                sum += a[i][j];
            printf ("Sum of %d-string equals %f\n",i,sum) ;
        }
    }
}
```



Data-Sharing Attribute Clauses

- Any variable in OpenMP is either **shared** or **private**
- Data-Sharing Attribute Clauses are used to explicitly define how variables should be scoped
- Data-Sharing Attribute Clauses include:
 - **shared,**
 - **private,**
 - **firstprivate,**
 - **lastprivate,**
 - **reduction,**
 - **default,**
 - **copyin,**
 - **copyprivate.**
- All the variables of the program are ***shared by default***



Data-Sharing Attribute Clauses

```
int sum = 0, a[NMAX], j;  
for (j=0; j < NMAX; j++)  
    sum += a[j];
```

- The code combines values into a single variable (sum)
 - there is a true dependence between loop iterations that can't be trivially removed

- Such situation is called a **Reduction** and it is frequently used.

REDUCTION Clause

Format: `reduction(op:list)`

- A private copy of each list variable is created for each thread of a team
- Local copies are initialized depending on the "op" (1 for "*" and 0 for "+")
- At the end of the reduction the final value of each of the local copies are combined with the original global value using the operator specified
- The variables in "list" must be shared in the enclosing parallel region



PI: PP with Worksharing and Reduction

```
#include <omp.h>
static long cntSteps=10000000; double step;

void main() {
    double x, pi, sum = 0.0;
    int i;
    step = 1.0/(double)cntSteps;
    #pragma omp parallel
    {
        double x;
        #pragma omp for reduction(+:sum)
        for (i=0; i<cntSteps; i++)    {
            x = (i + 0.5)*step;
            sum += 4.0/(1.0+ x*x) ;
        }
    }

    pi = sum*step;
    cout << "PI = "<<setw(15)<<setprecision(12)<<pi;
}
```

Data-Sharing Attribute Clauses

SHARED Clause

Format: `shared(list)`

- declares variables in its list to be shared among all threads in the team
- A shared variable exists in only one memory location and all threads can read or write to that address
- It is the programmer's responsibility to ensure that multiple threads properly access shared variables

DEFAULT Clause

Format: `default(shared|none)`

- allows the user to specify a default scope for all variables in the lexical extent of any parallel region
- Using **none** as a default requires that the programmer explicitly scope all variables
- Specific variables can be exempted from the default using the shared clause



Data-Sharing Attribute Clauses

PRIVATE Clause

Format: `private(list)`

- declares variables in its list to be private in each thread of the team
- Variables declared private should be uninitialized for each thread
- The original object referenced by the variable has an indeterminate value upon entry to the construct, and has an indeterminate value upon exit from the construct

FIRSTPRIVATE Clause

Format: `firstprivate(list)`

- private clause with automatic initialization of the variables in its list
- Listed variables are initialized according to the value of their original objects prior to entry into the parallel or work-sharing construct

LASTPRIVATE Clause

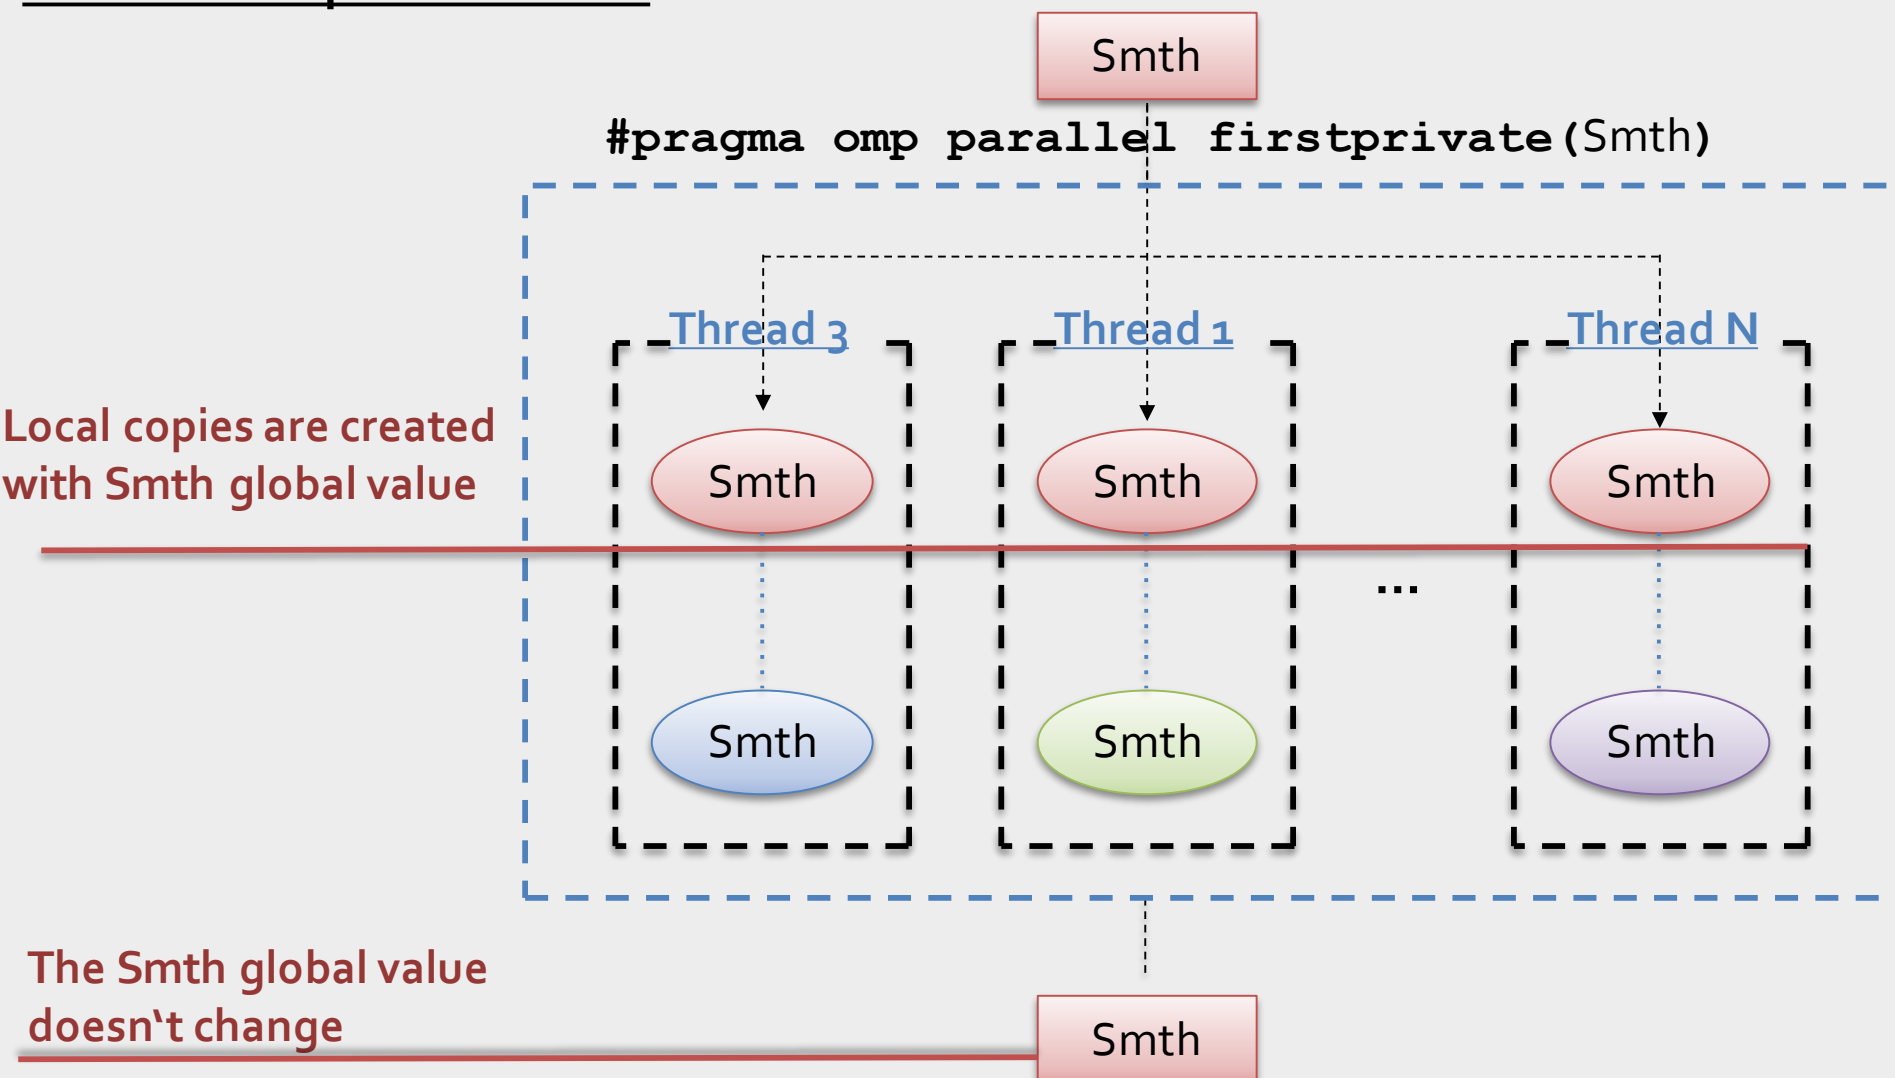
Format: `lastprivate(list)`

- private clause with a copy from the last (sequentially) loop iteration or section to the original variable object



Data-Sharing Attribute Clauses

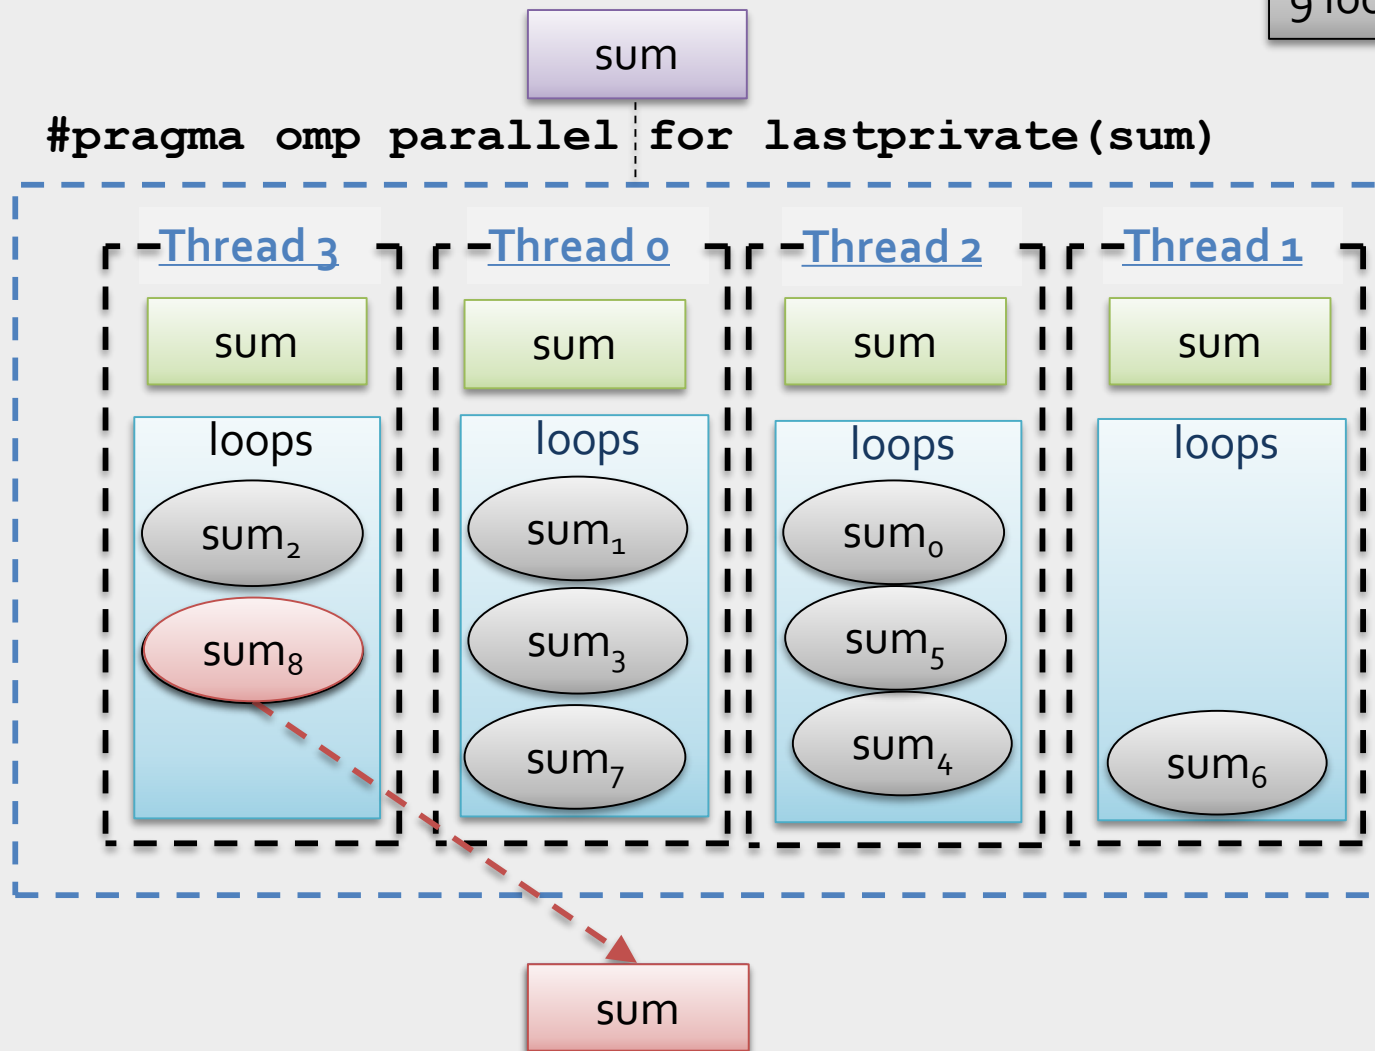
How the firstprivate works:



Data-Sharing Attribute Clauses

How the lastprivate works:

4 threads
9 loop iterations



Data-Sharing Attribute Clauses: Examples

```
int main() {  
    int b = 3;  
    b += 7;  
    printf ("b is %d\n", b);  
}
```



```
int b = 3;  
#pragma omp parallel firstprivate(b)  
{  
    b += 7;  
    printf ("b is %d\n", b);  
}
```



Synchronization Constructs in OpenMP

Sample:

```
int b = 3;
#pragma omp parallel
{
    b += 1;
    printf ("b is %d\n", b);
}
```



```
b is 6
b is 5
b is 9
b is 11
b is 7
```

- *Synchronization*: bringing one or more threads to a well defined and known point in their execution
- OpenMP includes:

High level synchronization



- barrier
- critical
- atomic
- ordered

Low level synchronization

- locks
- flush

Barrier

Mutual exclusion



Synchronization in OpenMP

BARRIER Directive

Format: `#pragma omp barrier`

- synchronizes all threads in the
- When a ***barrier*** directive is reached, a thread will wait at that point until all other threads have reached that barrier
- All threads then resume executing in parallel the code that follows the barrier
- All threads in a team (or none) must execute the ***barrier*** region



Synchronization in OpenMP

CRITICAL Directive

```
Format: #pragma omp critical [name]  
        structured_block
```

- specifies a region of code that must be executed by only one thread at a time
- If a thread is currently executing inside a ***critical*** region and another thread reaches that ***critical*** region and attempts to execute it, it will block until the first thread exits that ***critical*** region
- Names act as global identifiers. Different ***critical*** regions with the same name (or are unnamed) are treated as the same region
- It is illegal to branch into or out of a CRITICAL block



Synchronization in OpenMP

ATOMIC Directive

Format: `#pragma omp atomic`
`x <operator> = <expression>`

- specifies that a memory location must be updated atomically, rather than letting multiple threads attempt to write to it
- The directive applies only to a single, immediately following statement

Sample:

```
double A[NMAX];  
#pragma omp parallel for  
for (i = 0; i < NMAX; i++){  
    big_computations(A);  
  
    #pragma omp atomic  
    sum += A[i];  
}
```

The statement inside the atomic must be one of the following forms:

- `x binop= expr`
- `x++`
- `++x`
- `x—`
- `--x`

X is an lvalue of scalar type and binop is a non-overloaded built in operator.



Synchronization in OpenMP

Lock Functions

Initialize a
simple/nested lock

```
void omp_init_lock(omp_lock_t *lock);  
void omp_init_nest_lock(omp_nest_lock_t *lock);
```

Wait until a lock is
available

```
void omp_set_lock(omp_lock_t *lock);  
void omp_set_nest_lock(omp_nest_lock_t *lock);
```

Release a
simple/nested lock

```
void omp_unset_lock(omp_lock_t *lock);  
void omp_unset_nest_lock(omp_lock_t *lock);
```

Remove a
simple/nested lock

```
void omp_destroy_lock(omp_lock_t *lock);  
void omp_destroy_nest_lock(omp_nest_lock_t *lock);
```

Test a simple/nested
lock

```
int omp_test_lock(omp_lock_t *lock);  
int omp_test_nest_lock(omp_lock_t *lock);
```



Synchronization in OpenMP

Lock Functions. Example

```
int A[NMAX], sum = 0, i;
for (i = 0; i < NMAX; i++)
    A[i] = i;
omp_lock_t lock;
omp_init_lock(&lock);

#pragma omp parallel for
{
    for (i=0; i < NMAX; i++) {
        omp_set_lock (&lock);
        sum += A[i];
        omp_unset_lock (&lock);
    }
}
omp_destroy_lock (&lock);
```



HW: Tasks for report

- 2. Tell about:
 - Flush directive
 - Single directive
 - Master directive
 - Ordered directive
 - Sections Construct
- 3. Common Mistakes in OpenMP and How To Avoid Them (article)



Learning more about OpenMP

- OpenMP architecture review board URL, the “owner” of the OpenMP specification:
 - www.openmp.org
- OpenMP User’s Group (cOMPunity) URL:
 - www.compunity.org
- The OpenMP reference card:
 - <http://openmp.org/mp-documents/OpenMP3.1-CCard.pdf>



Getting Started. OpenMP vs. Visual Studio

To use OpenMP with Visual Studio, you need:

- Visual Studio 2005 Professional or better
- A multi processor or multi core system to see speed improvement
- An algorithm to parallelize

Visual Studio uses OpenMP standard 2.0 and is located in the vcomp.dll

To use OpenMP:

- Include `<omp.h>`
- Enable OpenMP compiler switch in Project Properties

